

Table of Contents

- [1. Preprocess Data](#)
 - [1.1 Data formating](#)
 - [1.2 Outliers](#)
 - [1.3 Missing values](#)
- [2. Transforms data](#)
 - [2.1 Aggregation](#)
 - [2.2 Logarithmization of variables](#)
 - [2.2 Handle with categorical features](#)
 - [2.3 Polynomial features](#)
- [3. LightGBM](#)

In [1]:

```
import math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import seaborn as sns
import lightgbm as lgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn import preprocessing
from sklearn.preprocessing import PolynomialFeatures
from sklearn.impute import SimpleImputer
from sklearn.model_selection import cross_val_score, cross_val_predict
from scipy.stats import shapiro
from scipy import stats
import os
print(os.listdir("house-prices-advanced-regression-techniques/"))

['test.csv', 'data_description.txt', 'train.csv', 'sample_submission.csv']
```

In [2]:

```
df_train = pd.read_csv("house-prices-advanced-regression-techniques/train.csv")
```

1. Preprocess Data

1.1 Data formating

'MSSubClass' - Identifies the type of dwelling involved in the sale. It is categorical feature despite Integer type. Before making one-hot encoding, use this feature to fill missed values in "Lot Frontage" and "BsmUnfSF", because seeing on boxplot we can see statistically significant difference between boxplots.

In [3]:

```
df_train.MSSubClass.isnull().sum()
```

Out[3]:

0

In [4]:

```
def fill_missing_with_mode_group(df, column):#разные/равные медианы в test & train
    groups = df["MSSubClass"].unique()
    for group in groups:
        median_in_group = df[df["MSSubClass"] == group][column].median()
        df.loc[df["MSSubClass"] == group, column] = df.loc[df["MSSubClass"] == group, column].fillna(median_in_group)
```

```
fill_missing_with_mode_group(df_train, 'LotFrontage')
```


```
fill_missing_with_mode_group(df_train, 'BsmtUnfSF')
```

In [5]:

```
for item in ['LotFrontage', 'BsmtUnfSF']:
    display(df_train.boxplot(item, by="MSSubClass", figsize=(5,5)))
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb7093dd850>
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb707345c40>
```



In [6]:

```
def formating_1(data):
    data = pd.concat([data, pd.get_dummies(data['MSSubClass'])], axis=1).drop('MSSu
    data.rename(columns={20:"1-STORY 1946 & NEWER ALL STYLES",
                          30:"1-STORY 1945 & OLDER",
                          40:"1-STORY W/FINISHED ATTIC ALL AGES",
                          45:"1-1/2 STORY - UNFINISHED ALL AGES",
                          50:"1-1/2 STORY FINISHED ALL AGES",
                          60:"2-STORY 1946 & NEWER",
                          70:"2-STORY 1945 & OLDER",
                          75:"2-1/2 STORY ALL AGES",
                          80:"SPLIT OR MULTI-LEVEL",
                          85:"SPLIT FOYER",
                          90:"DUPLEX - ALL STYLES AND AGES",
                          120:"1-STORY PUD (Planned Unit Development) - 1946 & NEWER",
                          150:"1-1/2 STORY PUD - ALL AGES",
                          160:"2-STORY PUD - 1946 & NEWER",
                          180:"PUD - MULTILEVEL - INCL SPLIT LEV/FOYER",
                          190:"2 FAMILY CONVERSION - ALL STYLES AND AGES"}},
                  inplace=True)
    return(data)

df_train = formating_1(df_train)
```

Features 'Condition1' & 'Condition2' have the same unique values.

In [7]:

```
pd.options.display.max_columns = None

def formating_2(data):
    data = pd.concat([data, pd.get_dummies(data['Condition1'])], axis=1).drop('Cond
    uniq_val = data['Condition2'].unique()
    for val in uniq_val:
        data.loc[data['Condition2']==val, val] = 1
    data.drop('Condition2', axis=1, inplace=True)
    return(data)

df_train = formating_2(df_train)

df_train[['Artery', 'Feedr', 'Norm', 'PosA', 'PosN', 'RRAe', 'RRAn', 'RRNe']].head(2)
```

Out[7]:

	Artery	Feedr	Norm	PosA	PosN	RRAe	RRAn	RRNe
0	0	0	1	0	0	0	0	0
1	0	1	1	0	0	0	0	0

In [8]:

```
df_train.shape
```

Out[8]:

```
(1460, 102)
```

Formatting qualitative features into numeric type

In [9]:

```
def convert_quality_into_numeric(data):
    data.ExterQual = data.ExterQual.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})
    data.ExterCond = data.ExterCond.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})
    data.BsmtQual = data.BsmtQual.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})#, np
    data.BsmtCond = data.BsmtCond.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})#, np
    data.BsmtExposure = data.BsmtExposure.map({'Gd':4, 'Av':3, 'Mn':2, 'No':1})#, np

    data.BsmtFinType1 = data.BsmtFinType1.map({'GLQ':6, 'ALQ':5, 'BLQ':4, 'Rec':3,
    data.BsmtFinType2 = data.BsmtFinType2.map({'GLQ':6, 'ALQ':5, 'BLQ':4, 'Rec':3,

    data.HeatingQC = data.HeatingQC.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})
    data.CentralAir = data.CentralAir.map({'N':0, 'Y':1})
    data.KitchenQual = data.KitchenQual.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})
    data.FireplaceQu = data.FireplaceQu.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})
    data.GarageFinish = data.GarageFinish.map({'Fin':3, 'RFn':2, 'Unf':1})#, np.nan
    data.GarageQual = data.GarageQual.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})
    data.GarageCond = data.GarageCond.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})
    data.PavedDrive = data.PavedDrive.map({'Y':3, 'P':2, 'N':1})
    data.PoolQC = data.PoolQC.map({'Ex':5, 'Gd':4, 'TA':3, 'Fa':2, 'Po':1})#, np.na

convert_quality_into_numeric(df_train)
```


1.2 Outliers

In [10]:

```
df_train.boxplot('SalePrice')
```

Out[10]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fb7093dd8e0>



In [11]:

```

outliers_threshold = df_train.SalePrice.quantile(0.75) + 1.5 * (df_train.SalePrice.
print(f"upper outliers threshold: {outliers_threshold},\n# outliers: {df_train.loc[

```

upper outliers threshold: 340037.5,
outliers: 61

Consider 61 observations, which are determined by outliers from the analysis of the boxplot.

In [12]:

```

def describe_outliers(threshold, MSZoning='none'):
    print('outliers statistics:')
    display(df_train.loc[(df_train.SalePrice > threshold) & (df_train.MSZoning == M
    print('Not outliers statistics: ')
    display(df_train.loc[(df_train.SalePrice <= threshold) & (df_train.MSZoning ==
    print('Observation with max price:')
    display(df_train.loc[(df_train.SalePrice == max(df_train.SalePrice)) & (df_train

```

Рассмотрим "выбросы" по группам.

In [13]:

```
df_train.groupby("MSZoning").SalePrice.agg(["count", "min", "median", "max"])
```

Out[13]:

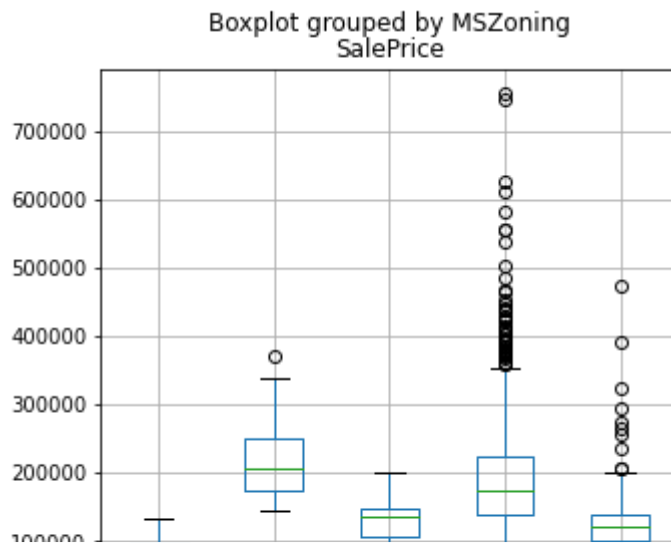
	count	min	median	max
MSZoning				
C (all)	10	34900	74700	133900
FV	65	144152	205950	370878
RH	16	76000	136500	200000
RL	1151	39300	174000	755000
RM	218	37900	120500	475000

In [16]:

```
df_train.boxplot(column="SalePrice", by="MSZoning", figsize=(5,5))
```

Out[16]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fb706d6bbe0>



In [17]:

```
def threshold_detector(MSZoning_type):
    IQR = df_train[df_train["MSZoning"] == MSZoning_type].SalePrice.quantile(.75) -
    thresh_value = df_train[df_train["MSZoning"] == MSZoning_type].SalePrice.quantile(.25)
    return(thresh_value)

rl_thresh = threshold_detector("RL")
rm_thresh = threshold_detector("RM")
```

MSZoning = "RL"

In [18]:

```
describe_outliers(rl_thresh, "RL")
```

outliers statistics:

	Id	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRem
mean	675.940000	89.080000	20324.000000	8.82000	5.120000	2000.460000	2003.0
std	399.008844	27.229066	29707.160623	0.84973	0.798979	15.718454	7.0
min	54.000000	42.000000	8089.000000	7.00000	2.000000	1934.000000	1965.0
25%	357.250000	69.000000	12393.000000	8.00000	5.000000	2003.000000	2003.0
50%	635.500000	88.000000	13792.000000	9.00000	5.000000	2006.000000	2006.0
75%	965.750000	105.000000	15097.750000	9.00000	5.000000	2008.000000	2008.0

Not outliers statistics:

	Id	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRen
mean	740.062670	73.620799	11174.348774	6.070845	5.554042	1974.326067	1984
std	421.034251	20.645238	8981.858533	1.277951	1.056596	25.824160	19
min	1.000000	24.000000	2268.000000	1.000000	1.000000	1875.000000	1950
25%	373.000000	61.000000	8400.000000	5.000000	5.000000	1958.000000	1968
50%	754.000000	75.000000	9880.000000	6.000000	5.000000	1975.000000	1992
75%	1103.000000	80.000000	11900.000000	7.000000	6.000000	1999.000000	2003

Observation with max price:

	Id	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasV
mean	692.0	104.0	21535.0	10.0	6.0	1994.0	1995.0	

In the case of "Outliers", there is a large std in the LotArea variable (area in sq.m.).

In [19]:

```
rl_outliers = df_train.loc[(df_train.SalePrice > rl_thresh) & (df_train.MSZoning ==
#rl_outliers.head(15)

# OverallQual - [0, 10]
# OverallCond - [0, 10]
```

OverallCond is the same for all non-outliers OverallQual for all > 7.0 [3 quantile not outliers] SalePrice is noticeably more than 75% of non-outliers.

It was decided to determine observations with a small area (≤ 13792) as outliers Let's remove outliers.

In [20]:

```
display(rl_outliers[rl_outliers["LotArea"] <= 13792].describe().loc[["count", "mean"]
rl_outliers = rl_outliers[rl_outliers["LotArea"] <= 13792]
df_train.drop(axis = 0, index=rl_outliers.index, inplace=True)
print(f'number deleted observations: {rl_outliers.shape[0]}')
```

	LotArea	OverallQual	OverallCond	SalePrice
count	25.000000	25.00000	25.000000	25.000000
mean	11974.080000	8.84000	4.960000	408791.680000
std	1545.240648	0.85049	0.675771	54260.760228

number deleted observations: 25

In [21]:

```
#describe_outliers(rm_thresh, "RM")
```

All observations have OverallQual and OverallCond about 75% not outliers. Delete observations, with LotArea <8520

In [22]:

```
rm_outliers = df_train.loc[(df_train.SalePrice > rm_thresh) & (df_train.MSZoning ==
rm_outliers = rm_outliers[rm_outliers["LotArea"] <= 8520]
df_train.drop(axis = 0, index=rm_outliers.index, inplace=True)
print(f'number deleted observations: {rm_outliers.shape[0]}')
```

number deleted observations: 5

1.3 Missing values

In [23]:

```
def missing_ration(data, top=10):
    df_na = round((data.isnull().sum() / len(data)) * 100,2)
    df_na = df_na.drop(df_na[df_na == 0].index).sort_values(ascending=False)
    missing_data = pd.DataFrame(
        {'Missing Ratio' :df_na}
    )
    print(missing_data)

def fill_missing(df, cols, val):
    """ Replace with the supplied val """
    for col in cols:
        df[col] = df[col].fillna(val)

def fill_missing_with_mode(df, cols):
    """ Replace with the mode """
    for col in cols:
        df[col] = df[col].fillna(df[col].mode()[0])

print("\n Train")
missing_ration(df_train)
```

Train	Missing Ratio
PoolQC	99.51
MiscFeature	96.22
Alley	93.64
Fence	80.35
FireplaceQu	48.18
GarageCond	5.66
GarageQual	5.66
GarageFinish	5.66
GarageYrBlt	5.66
GarageType	5.66
BsmtFinType2	2.66
BsmtExposure	2.66
BsmtFinType1	2.59
BsmtCond	2.59
BsmtQual	2.59
MasVnrArea	0.56
MasVnrType	0.56
Electrical	0.07

Replace with "Not exist"

Reading 'data_description.txt' notice that some there are some variables, which contains NA/None and it has meaning. Replace missing values with "Not exist"

In [24]:

```
fill_missing(df_train, ["MiscFeature", "Alley", "Fence",
                        "GarageType", "MasVnrType"], "Not exist")
```

Replace with 0

Missed values of Quantities features replace with 0.

Replace with mode

Missed values of categorical variables replace with mode, because not many missing values.

In [25]:

```
fill_missing(df_train, ['GarageYrBlt', 'MasVnrArea', 'BsmtHalfBath', 'BsmtFullBath',  
                        'BsmtFinSF1', 'GarageArea', 'TotalBsmtSF', 'GarageCars', 'B  
                        "PoolQC", "FireplaceQu", "GarageFinish", "GarageQual", "Gara  
                        "BsmtCond", "BsmtExposure", "BsmtFinType1", "BsmtFinType2"],  
  
fill_missing_with_mode(df_train, ['Electrical', 'MSZoning', 'Functional', 'Utilitie  
'KitchenQual', 'SaleType'])
```

2. Transforms data

2.1 Aggregation

In [26]:

```
df_train['TotalSF'] = df_train['TotalBsmtSF'] + df_train['1stFlrSF'] + df_train['2n
```

In [27]:

```
df_train["time_before_remodelation"] = df_train.YearRemodAdd - df_train.YearBuilt
```

2.2 Logarithmization of variables

Taking the logarithm of some variables allows their distribution to be closer to normal ('loglist' tuple was created by seeing histograms and shapiro test)

In [28]:

```
def add_logs(df, cols):  
    # add Log transformed columns  
    for c in cols:  
        df[c] = np.log(df[c] + 1)  
    return(df)  
  
loglist = ('LotFrontage',  
           'LotArea', 'BsmtUnfSF', 'GrLivArea',  
           'BsmtFullBath', 'BsmtHalfBath', 'TotalSF')  
  
df_train = add_logs(df_train, loglist)
```

In [29]:

```
df_train["SalePrice"] = np.log1p(df_train["SalePrice"])
```

2.3 Handle with categorical features

In [30]:

```
def fix_missing_cols(in_train, in_test):
    missing_cols = set(in_train.columns) - set(in_test.columns)
    # Add a missing column in test set with default value equal to 0
    for c in missing_cols:
        in_test[c] = 0
    # Ensure the order of column in the test set is in the same order than in train
    in_test = in_test[in_train.columns]
    return(in_test)

def dummy_encode(in_df_train):
    df_train = in_df_train
    categorical_feats = [
        f for f in df_train.columns if df_train[f].dtype == 'object'
    ]
    print(categorical_feats)
    for f_ in categorical_feats:
        prefix = f_
        df_train = pd.concat([df_train, pd.get_dummies(df_train[f_], prefix=prefix)])

    return(df_train)
```

In [31]:

```
df_train = dummy_encode(df_train)
df_train_here = df_train
print("Shape train: %s, test: %s" % (df_train.shape))
```

```
['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities',
 'LotConfig', 'LandSlope', 'Neighborhood', 'BldgType', 'HouseStyle',
 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
 'Foundation', 'Heating', 'Electrical', 'Functional', 'GarageType',
 'Fence', 'MiscFeature', 'SaleType', 'SaleCondition']
Shape train: 1430, test: 251
```

2.4 Polynomial features

Make polynomial features of variables, which have high correlation coefficient with Dependent Variable

In [32]:

```
def load_poly_features(data, cols):
    print('Loading polynomial features..')
    # Make a new dataframe for polynomial features
    poly_features = data[cols]

    # imputer for handling missing values
    imputer = SimpleImputer(strategy='median')

    # Need to impute missing values
    poly_features = imputer.fit_transform(poly_features)

    # Create the polynomial object with specified degree
    poly_transformer = PolynomialFeatures(degree=3)
    # Train the polynomial features
    poly_features = poly_transformer.fit_transform(poly_features)

    print('Polynomial Features shape: %s' % str(poly_features.shape))

    df_poly_features = pd.DataFrame(poly_features,
                                     columns=poly_transformer.get_feature_names(cols))

    df_poly_features['Id'] = data['Id']

    print('Loaded polynomial features')
    return(df_poly_features)
```

In [33]:

```
corr_list = abs(df_train.corr().SalePrice).sort_values(ascending=False)
correlated_cols = corr_list[corr_list > 0.7][1:]
print(correlated_cols)
correlated_cols = correlated_cols.index
```

```
OverallQual    0.803943
TotalSF        0.797013
GrLivArea      0.722385
Name: SalePrice, dtype: float64
```

In [34]:

```
df_train_poly = load_poly_features(df_train, cols=correlated_cols)
print("Shape train: %s, test: %s" % (df_train_poly.shape))

df_train = df_train.merge(right=df_train_poly.reset_index(), how='left', on='Id')

print("Shape train: %s, test: %s" % (df_train.shape))
```

```
Loading polynomial features..
Polynomial Features shape: (1430, 20)
Loaded polynomial features
Shape train: 1430, test: 21
Shape train: 1430, test: 272
```

3. Light GBM

Here for the sake of higher score on kaggle, **I commented out some lines, which ordinary used**. It allowed me went from 21% to 17% in leaderborad.

In [35]:

```
y = df_train["SalePrice"]
df_train_id = df_train.pop('Id')
df_train.drop(["SalePrice"], axis=1, inplace=True)
print("Shape train: %s, test: %s" % (df_train.shape))
```

Shape train: 1430, test: 270

In [36]:

```
X_train = df_train
y_train = y
#X_train, X_test, y_train, y_test = train_test_split( df_train, y, test_size=0.1, r
```

The best parameters was found by using *GridSearch*

In []:

```
'''
from sklearn.model_selection import GridSearchCV
hyper_params = {
    'task': 'train',
    'boosting_type': 'gbdt',
    'objective': 'regression',
    'metric': ['l2'],
    'first_metric_only': True,
    'bagging_freq': 10,
    'verbose': 0,
    "num_leaves": 128,
    "num_iterations": 3000
}

gbm = lgb.LGBMRegressor(**hyper_params)
parameters = {
    'learning_rate': [0.001],
    'feature_fraction': [0.5, 0.7],
    'bagging_fraction': [0.6, 0.75],
    "max_depth": [4, 5, 6],
    'bagging_freq': [10],
    "max_bin": [300, 400, 500],
    "n_estimators": [200, 300, 400]
}

grid = GridSearchCV(estimator=gbm, param_grid=parameters,
                    cv=5, verbose=500, n_jobs=-1, refit=True)
grid.fit(X_train, y_train)
print(f"{grid.best_estimator_}\n{grid.best_params_}\n{grid.best_score_}")
'''
```

In [37]:

```
best_parameters = {
    'task': 'train',
    'boosting_type': 'gbdt',
    'objective': 'regression',
    'metric': ['l2', 'huber'],
    'first_metric_only': True,
    'learning_rate': 0.005,
    'feature_fraction': 0.5,
    'bagging_fraction': 0.7,
    'bagging_freq': 10,
    'verbose': 0,
    "max_depth": 6,
    "num_leaves": 128,
    "max_bin": 512,
    "num_iterations": 3000,
    "n_estimators": 50
}

gbm = lgb.LGBMRegressor(**best_parameters)
```

In [38]:

```
gbm.fit(X_train, y_train,
#       eval_set=(X_test, y_test),
#       eval_metric='l2',
#       early_stopping_rounds=1000,
#       verbose = 500)
```

```
/home/user/anaconda3/lib/python3.8/site-packages/lightgbm/engine.py:14
8: UserWarning: Found `num_iterations` in params. Will use it instead
of argument
  warnings.warn("Found `{}` in params. Will use it instead of argumen
t".format(alias))
```

Out[38]:

```
LGBMRegressor(bagging_fraction=0.7, bagging_freq=10, feature_fraction=
0.5,
               first_metric_only=True, learning_rate=0.005, max_bin=51
2,
               max_depth=6, metric=['l2', 'huber'], n_estimators=50,
               num_iterations=3000, num_leaves=128, objective='regressi
on',
               task='train', verbose=0)
```

Prediction

In [39]:

```
y_pred = gbm.predict(X_train, num_iteration=gbm.best_iteration_)
print('The rmse of prediction is:', round(mean_squared_error(y_pred, y_train) ** 0.5, 4))

The rmse of prediction is: 0.05951
```

Results

In [40]:

```
def test_pipeline(data):
    fill_missing_with_mode_group(data, 'LotFrontage')
    fill_missing_with_mode_group(data, 'BsmtUnfSF')
    data = formating_1(data)
    data = formating_2(data)
    convert_quality_into_numeric(data)
    fill_missing(data, ["MiscFeature", "Alley", "Fence",
                        "GarageType", "MasVnrType"], "Not exist")
    fill_missing(data, ['GarageYrBlt', 'MasVnrArea', 'BsmtHalfBath', 'BsmtFullBath',
                        'BsmtFinSF1', 'GarageArea', 'TotalBsmtSF', 'GarageCars',
                        "PoolQC", "FireplaceQu", "GarageFinish", "GarageQual", "
                        "BsmtCond", "BsmtExposure", "BsmtFinType1", "BsmtFinType2",
                        'Electrical', 'MSZoning', 'Functional', 'Utilities',
                        'KitchenQual', 'SaleType'])
    data['TotalSF'] = data['TotalBsmtSF'] + data['1stFlrSF'] + data['2ndFlrSF']
    data["time_before_remodelation"] = data.YearRemodAdd - data.YearBuilt
    data = add_logs(data, loglist)
    data = dummy_encode(data)
    data = fix_missing_cols(df_train_here, data)
    data_poly = load_poly_features(data, cols=correlated_cols)
    data = data.merge(right=data_poly.reset_index(), how='left', on='Id')
    data.drop("SalePrice", axis=1, inplace=True)
    return(data)
```

In [41]:

```
df_test = test_pipeline(pd.read_csv("house-prices-advanced-regression-techniques/test.csv"))
df_test_id = df_test.pop('Id')
test_pred = np.expm1(gbm.predict(df_test, num_iteration=gbm.best_iteration_))

df_test["SalePrice"] = test_pred
df_test["Id"] = df_test_id
df_test.to_csv("results.csv", columns=["Id", "SalePrice"], index=False)
```

```
['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities',
 'LotConfig', 'LandSlope', 'Neighborhood', 'BldgType', 'HouseStyle',
 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
 'Foundation', 'Heating', 'Electrical', 'Functional', 'GarageType',
 'Fence', 'MiscFeature', 'SaleType', 'SaleCondition']
```

Loading polynomial features..

Polynomial Features shape: (1459, 20)

Loaded polynomial features

Score: 0.12301 (17% [825/4675])

<https://www.kaggle.com/konstantinlp/competitions> (<https://www.kaggle.com/konstantinlp/competitions>)